

Introduction to Computer Systems, Spring 2012

Lab Assignment L3: The Buffer Bomb

Assigned: April 5 Due: April 16, 15:59

Introduction

This assignment helps you develop a detailed understanding of the calling stack organization on an IA32(x86) processor. It involves applying a series of *buffer overflow attacks* on an executable file `bufbomb` in the lab directory.

Note: In this lab, you will gain firsthand experience with one of the methods commonly used to exploit security weaknesses in operating systems and network servers. Our purpose is to help you learn about the runtime operation of programs and to understand the nature of this form of security weakness so that you can avoid it when you write system code. We do not condone the use of these or any other form of attack to gain unauthorized access to any system resources. There are criminal statutes governing such activities.

Logistics

You should work alone(only one people) in solving the problems for this assignment. The only “hand-in” will be an automated logging of your successful attacks. Any clarifications and revisions to the assignment will be posted on the course Web page or fudan BBS.

Hand Out Instructions

The lab handout is now available on the course web page (<http://ipads.se.sjtu.edu.cn/courses/ics/schedule.shtml>), where you can also obtain your lab status information.

Start by copying `buflab-handout.tar` to a (protected) directory in which you plan to do your work. Then give the command “`tar xvf buflab-handout.tar`”. This will cause a number of files to be unpacked in the directory:

MAKECOOKIE: Generates a “cookie” based on your team name.

BUFBOMB: The code you will attack.

SENDSTRING: A utility to help convert between string formats.

All of these programs are compiled to run on Linux machines.

In the following instructions, we will assume that you have copied the three programs to a protected local directory, and that you are executing them in that local directory.

Stack Randomization

As mentioned in ICS book section 3.12.1, Stack Randomization has widely adopted by current linux distributions (e.g., debian). This mechanism would affect the verification of buflab. The most simple solution is to prohibit it by following command.

```
unix> echo 0 > /proc/sys/kernel/randomize_va_space
```

But the setting is temporal, so you need to run above command everytime after you start or reboot your machine. You could use the following command to check the current status of machine, and the output should be zero.

```
unix> cat /proc/sys/kernel/randomize_va_space
```

Team Name and Cookie

It is required that each team consists of exactly one student. So the team name is exactly your **ICS SVN id** (e.g. ics5110379076).

A *cookie* is a string of eight hexadecimal digits that is (with high probability) uniquely generated for your team. You can generate your cookie with the `makecookie` program giving your team name as the argument. For example:

```
unix> ./makecookie ics5110379076
0x4e5601e3
```

In your four buffer attacks, your objective will be to make your cookie show up in places where it ordinarily would not.

The BUFBOMB Program

The BUFBOMB program reads a string from standard input with a function `getbuf` having the following C code:

```
1 int getbuf()  
2 {  
3     char buf[12];
```

```

4     Gets(buf);
5     return 1;
6 }

```

The function `Gets` is similar to the standard library function `gets`—it reads a string from standard input (terminated by ‘\n’ or end-of-file) and stores it (along with a null terminator) at the specified destination. In this code, the destination is an array `buf` having sufficient space for 12 characters.

Neither `Gets` nor `gets` has any way to determine whether there is enough space at the destination to store the entire string. Instead, they simply copy the entire string, possibly overrunning the bounds of the storage allocated at the destination.

If the string typed by the user to `getbuf` is no more than 11 characters long, it is clear that `getbuf` will return 1, as shown by the following execution example:

```

unix> ./bufbomb -t ics5110379076
Type string: ics5110379076
Dud: getbuf returned 0x1

```

Typically an error occurs if we type a longer string:

```

unix> ./bufbomb
Type string: This string is too long
Ouch!: You caused a segmentation fault!

```

As the error message indicates, overrunning the buffer typically causes the program state to be corrupted, leading to a memory access error. Your task is to be more clever with the strings you feed `BUFBOMB` so that it does more interesting things. These are called *exploit* strings.

`BUFBOMB` takes several different command line arguments:

-t *TEAM*: Operate the bomb for the indicated team. You should always provide this argument for several reasons:

- It is required to log your successful attacks.
- `BUFBOMB` determines the cookie you will be using based on your team name, just as does the program `MAKECOOKIE`.
- We have built features into `BUFBOMB` so that some of the key stack addresses you will need to use depend on your team’s cookie.

-t *team*: Specify team name

-s: Submit solution to server

-h: Print list of possible command line arguments

Your exploit strings will typically contain byte values that do not correspond to the ASCII values for printing characters. The program SENDSTRING can help you generate these *raw* strings. It takes as input a *hex-formatted* string. In this format, each byte value is represented by two hex digits. For example, the string “012345” could be entered in hex format as “30 31 32 33 34 35.” (Recall that the ASCII code for decimal digit x is $0x3x$.) Non-hex digit characters are ignored, including the blanks in the example shown.

If you generate a hex-formatted exploit string in the file `exploit.txt`, you can apply the raw string to BUFBOMB in several different ways:

1. You can set up a series of pipes to pass the string through SENDSTRING.

```
unix> cat exploit.txt | ./sendstring | ./bufbomb -t ics5110379076
```

2. You can store the raw string in a file and use I/O redirection to supply it to BUFBOMB:

```
unix> ./sendstring < exploit.txt > exploit-raw.txt
unix> ./bufbomb -t ics5110379076 < exploit-raw.txt
```

This approach can also be used when running BUFBOMB from within GDB:

```
unix> gdb bufbomb (gdb) run -t ics5110379076 < exploit-raw.txt
```

One important point: your exploit string must not contain byte value `0x0A` at any intermediate position, since this is the ASCII code for newline (`'\n'`). When `Gets` encounters this byte, it will assume you intended to terminate the string. SENDSTRING will warn you if it encounters this byte value.

When you correctly solve one of the levels, BUFBOMB will automatically send an email notification to our grading server by using argument `-s`.

```
unix> cat exploit.txt | ./sendstring | ./bufbomb -t ics5110379076 -s
```

The server will test your exploit string to make sure it really works, and it will update the lab web page indicating that your team (listed by cookie) has completed this level.

Unlike the bomb lab, there is no penalty for making mistakes in this lab. Feel free to fire away at BUFBOMB with any string you like.

Level 0: Candle (10 pts)

The function `getbuf` is called within BUFBOMB by a function `test` having the following C code:

```
1 void test()
2 {
3     int val;
4     volatile int local = 0xdeadbeef;
5     entry_check(3); /* Make sure entered this function properly */
6     val = getbuf();
```

```

7      /* Check for corrupted stack */
8      if (local != 0xdeadbeef) {
9          printf("Sabotaged!: the stack has been corrupted\n");
10     }
11     else if (val == cookie) {
12         printf("Boom!: getbuf returned 0x%x\n", val);
13         validate(3);
14     }
15     else {
16         printf("Dud: getbuf returned 0x%x\n", val);
17     }
18 }

```

When `getbuf` executes its return statement (line 5 of `getbuf`), the program ordinarily resumes execution within function `test` (at line 8 of this function). Within the file `bufbomb`, there is a function `smoke` having the following C code:

```

void smoke()
{
    entry_check(0); /* Make sure entered this function properly */
    printf("Smoke!: You called smoke()\n");
    validate(0);
    exit(0);
}

```

Your task is to get `BUFBOMB` to execute the code for `smoke` when `getbuf` executes its return statement, rather than returning to `test`. You can do this by supplying an exploit string that overwrites the stored return pointer in the stack frame for `getbuf` with the address of the first instruction in `smoke`. Note that your exploit string may also corrupt other parts of the stack state, but this will not cause a problem, since `smoke` causes the program to exit directly.

Some Advice:

- All the information you need to devise your exploit string for this level can be determined by examining a disassembled version of `BUFBOMB`.
- Be careful about byte ordering.
- You might want to use GDB to step the program through the last few instructions of `getbuf` to make sure it is doing the right thing.
- The placement of `buf` within the stack frame for `getbuf` depends on which version of GCC was used to compile `bufbomb`. You will need to pad the beginning of your exploit string with the proper number of bytes to overwrite the return pointer. The values of these bytes can be arbitrary.

Level 1: Sparkler (20 pts)

Within the file `bufbomb` there is also a function `fizz` having the following C code:

```

void fizz(int val)
{
    entry_check(1); /* Make sure entered this function properly */
    if (val == cookie) {
        printf("Fizz!: You called fizz(0x%x)\n", val);
        validate(1);
    } else
        printf("Misfire: You called fizz(0x%x)\n", val);
    exit(0);
}

```

Similar to Level 0, your task is to get BUFBOMB to execute the code for `fizz` rather than returning to `test`. In this case, however, you must make it appear to `fizz` as if you have passed your cookie as its argument. You can do this by encoding your cookie in the appropriate place within your exploit string.

Some Advice:

- Note that the program won't really call `fizz`—it will simply execute its code. This has important implications for where on the stack you want to place your cookie.

Level 2: Firecracker (30 pts)

A much more sophisticated form of buffer attack involves supplying a string that encodes actual machine instructions. The exploit string then overwrites the return pointer with the starting address of these instructions. When the calling function (in this case `getbuf`) executes its `ret` instruction, the program will start executing the instructions on the stack rather than returning. With this form of attack, you can get the program to do almost anything. The code you place on the stack is called the *exploit* code. This style of attack is tricky, though, because you must get machine code onto the stack and set the return pointer to the start of this code.

Within the file `bufbomb` there is a function `bang` having the following C code:

```

int global_value = 0;

void bang(int val)
{
    entry_check(2); /* Make sure entered this function properly */
    if (global_value == cookie) {
        printf("Bang!: You set global_value to 0x%x\n", global_value);
        validate(2);
    } else
        printf("Misfire: global_value = 0x%x\n", global_value);
    exit(0);
}

```

Similar to Levels 0 and 1, your task is to get BUFBOMB to execute the code for `bang` rather than returning to `test`. Before this, however, you must set global variable `global_value` to your team's cookie. Your

exploit code should set `global_value`, push the address of `bang` on the stack, and then execute a `ret` instruction to cause a jump to the code for `bang`.

Some Advice:

- You can use GDB to get the information you need to construct your exploit string. Set a breakpoint within `getbuf` and run to this breakpoint. Determine parameters such as the address of `global_value` and the location of the buffer.
- Determining the byte encoding of instruction sequences by hand is tedious and prone to errors. You can let tools do all of the work by writing an assembly code file containing the instructions and data you want to put on the stack. Assemble this file with `GCC` and disassemble it with `OBJDUMP`. You should be able to get the exact byte sequence that you will type at the prompt. (A brief example of how to do this is included at the end of this writeup.)
- Keep in mind that your exploit string depends on your machine, your compiler, and even your team's cookie. Do all of your work on a Fish machine, and make sure you include the proper team name on the command line to `BUFBOMB`.
- Our solution requires 16 bytes of exploit code. Fortunately, there is sufficient space on the stack, because we can overwrite the stored value of `%ebp`. This stack corruption will not cause any problems, since `bang` causes the program to exit directly.
- Watch your use of address modes when writing assembly code. Note that `movl $0x4, %eax` moves the *value* `0x00000004` into register `%eax`; whereas `movl 0x4, %eax` moves the value *at* memory location `0x00000004` into `%eax`. Since that memory location is usually undefined, the second instruction will cause a segfault!
- Do not attempt to use either a `jmp` or a `call` instruction to jump to the code for `bang`. These instructions use PC-relative addressing, which is very tricky to set up correctly. Instead, push an address on the stack and use the `ret` instruction.

Level 3: Dynamite (40 pts)

Our preceding attacks have all caused the program to jump to the code for some other function, which then causes the program to exit. As a result, it was acceptable to use exploit strings that corrupt the stack, overwriting the saved value of register `%ebp` and the return pointer.

The most sophisticated form of buffer overflow attack causes the program to execute some exploit code that patches up the stack and makes the program return to the original calling function (`test` in this case). The calling function is oblivious to the attack. This style of attack is tricky, though, since you must: 1) get machine code onto the stack, 2) set the return pointer to the start of this code, and 3) undo the corruptions made to the stack state.

Your job for this level is to supply an exploit string that will cause `getbuf` to return your cookie back to `test`, rather than the value 1. You can see in the code for `test` that this will cause the program to go

“Boom!” Your exploit code should set your cookie as the return value, restore any corrupted state, push the correct return location on the stack, and execute a `ret` instruction to really return to `test`.

Some Advice:

- In order to overwrite the return pointer, you must also overwrite the saved value of `%ebp`. However, it is important that this value is correctly restored before you return to `test`. You can do this by either 1) making sure that your exploit string contains the correct value of the saved `%ebp` in the correct position, so that it never gets corrupted, or 2) restore the correct value as part of your exploit code. You’ll see that the code for `test` has some explicit tests to check for a corrupted stack.
- You can use GDB to get the information you need to construct your exploit string. Set a breakpoint within `getbuf` and run to this breakpoint. Determine parameters such as the saved return address and the saved value of `%ebp`.
- Again, let tools such as GCC and OBJDUMP do all of the work of generating a byte encoding of the instructions.
- Keep in mind that your exploit string depends on your machine, your compiler, and even your team’s cookie. Do all of your work on a Fish machine, and make sure you include the proper team name on the command line to BUFBOMB.

Once you complete this level, pause to reflect on what you have accomplished. You caused a program to execute machine code of your own design. You have done so in a sufficiently stealthy way that the program did not realize that anything was amiss.

Logistical Notes

Hand in occurs automatically whenever you correctly solve a level. The program sends email to our grading server containing your team name (be sure to set the “-t” command line flag properly) and your exploit string to the grading server. You will be informed of this by BUFBOMB. Upon receiving the email, the server will validate your string and update the lab web page. You should check this page a few minutes after your submission to make sure your string has been validated. [If you really solved the level, your string *should* be valid.]

Note that each level is graded individually. You do not need to do them in the specified order, but you will get credit only for the levels for which the server receives a valid message.

Have fun!

Generating Byte Codes

Using GCC as an assembler and OBJDUMP as a disassembler makes it convenient to generate the byte codes for instruction sequences. For example, suppose we write a file `example.s` containing the following assembly code:


```
# Example of hand-generated assembly code
    pushl $0x89abcdef      # Push value onto stack
    addl $17,%eax          # Add 17 to %eax
    .align 4               # Following will be aligned on multiple of 4
    .long 0xfedcba98       # A 4-byte constant
    .long 0x00000000       # Padding
```

The code can contain a mixture of instructions and data. Anything to the right of a ‘#’ character is a comment. We have added an extra word of all 0s to work around a shortcoming in OBJDUMP to be described shortly.

We can now assemble and disassemble this file:

```
unix> gcc -c example.s
unix> objdump -d example.o > example.d
```

The generated file `example.d` contains the following lines

```
0: 68 ef cd ab 89      push    $0x89abcdef
5: 83 c0 11            add     $0x11,%eax
8: 98                  cwtl
9: ba dc fe 00 00     mov     $0xfedc,%edx  Objdump tries to interpret
                        these as instructions
```

Each line shows a single instruction. The number on the left indicates the starting address (starting with 0), while the hex digits after the ‘:’ character indicate the byte codes for the instruction. Thus, we can see that the instruction `pushl $0x89ABCDEF` has hex-formatted byte code `68 ef cd ab 89`.

Starting at address 8, the disassembler gets confused. It tries to interpret the bytes in the file `example.o` as instructions, but these bytes actually correspond to data. Note, however, that if we read off the 4 bytes starting at address 8 we get: `98 ba dc fe`. This is a byte-reversed version of the data word `0xFEDCBA98`. This byte reversal represents the proper way to supply the bytes as a string, since a little endian machine lists the least significant byte first. Note also that it only generated two of the four bytes at the end with value `00`. Had we not added this padding, OBJDUMP gets even more confused and does not emit all of the bytes we want.

Finally, we can read off the byte sequence for our code (omitting the final 0’s) as:

```
68 ef cd ab 89 83 c0 11 98 ba dc fe
```